
Representation Regularization and Theoretical Mechanics of Joint-Embedding Predictive Architectures for Multi-Modal Brain MRI: A Comprehensive Analysis of Label Efficiency, Non-Collapse Dynamics, and Out-of-Distribution Generalization

Han Riman *

Master’s Thesis Laboratory
hanriman@master.thesis

Abstract

Joint-Embedding Predictive Architectures (JEPA) represent a paradigm shift in self-supervised visual representation learning, operating entirely in latent embedding space without pixel-level image reconstruction. Standard I-JEPA prevents latent representation collapse through an Exponential Moving Average (EMA) momentum update of a teacher target encoder. Recent theoretical formulations—specifically SIGReg (Sketched Isotropic Gaussian Regularization in LeJEPA) and VISReg (Variance-Invariance-Sketching Regularization)—introduce analytical feature regularization penalties that enable *heuristic-free* single-encoder JEPA training without momentum teachers. In this paper, we present a mathematically rigorous and comprehensive evaluation of I-JEPA, SigReg JEPA, VisReg JEPA, a 2D Residual UNet baseline, and a state-of-the-art 2D nnU-Net baseline with deep supervision on 2D multi-modal Glioma MRI (BraTS GLI). We detail the mathematical formulations and architectural micro-designs across all models, including internal residual skip connections, multi-scale encoder-to-decoder concatenation skip connections, pure bottleneck transposed convolutional decoding, stochastic patch block masking with batch-uniform collation, Random Modality Dropout with guaranteed active channel fallbacks, hyperspherical projection sketching, and gradient-clipped optimization schedules. We evaluate latent collapse metrics (Effective Feature Rank Rank_{eff} , Average Pairwise Cosine Similarity S_{cosine}), downstream segmentation metrics (Dice Similarity Coefficient, Intersection over Union, 95th Percentile Hausdorff Distance HD95), and out-of-distribution transfer across three clinical regimes: (1) Full-Data Benchmark, (2) Low-Data Label Efficiency (1% to 100% labels), and (3) Out-of-Distribution (OOD) Scanner Shift and Zero-Shot Cross-Pathology Transfer on Meningioma MRI (BraTS-MEN-RT). Our empirical findings establish that under extreme label scarcity (1% labels / 63 training slices), SigReg JEPA achieves a top downstream Test Dice of 0.5933, outperforming fully supervised nnU-Net (0.3847) by +20.86% absolute Dice (+54.2% relative improvement) while executing fine-tuning $1.9\times$ faster per epoch (0.75 s/epoch vs 1.42 s/epoch). Furthermore, under zero-padded missing modality shifts, self-supervised JEPAs preserve zero-shot transferability (0.0700–0.0929 Dice) whereas standard UNet completely collapses (0.0151 Dice), proving that analytical variance-covariance regularization builds

*Corresponding author. Master’s Thesis Research Laboratory.

modality-invariant, high-rank representations while eliminating momentum teacher overhead.

1 Introduction and Theoretical Framework

Automated segmentation of brain gliomas from multi-modal Magnetic Resonance Imaging (MRI)—comprising T1-weighted (T1), post-contrast T1-weighted (T1c), T2-weighted (T2), and Fluid Attenuated Inversion Recovery (FLAIR) sequences—is essential for clinical neuro-oncology diagnosis, surgical boundary planning, and longitudinal monitoring [Baid et al., 2024]. Supervised deep convolutional architectures, such as Residual UNet [Ronneberger et al., 2015] and nnU-Net [Isensee et al., 2021], achieve high accuracy when extensive pixel-wise ground-truth annotations are available. However, expert manual voxel delineation is labor-intensive, scarce, and susceptible to inter-observer variability.

Self-Supervised Learning (SSL) addresses this bottleneck by pre-training representation encoders on unannotated MRI volumes prior to downstream task fine-tuning. Unlike Masked Autoencoders (MAE) [He et al., 2022] that reconstruct high-frequency pixel intensities in pixel space, Joint-Embedding Predictive Architectures (JEPA) [Assran et al., 2023] operate exclusively in latent embedding space. JEPA models predict target patch representations from context patch representations, forcing the encoder to learn abstract semantic structures rather than low-level noise.

A central challenge in energy-based self-supervised architectures is preventing *latent representation collapse*—a trivial state where the encoder maps all input patches to a constant vector $z = c \in \mathbb{R}^D$, yielding zero prediction error $\mathcal{L} = 0$ without learning meaningful semantics.

Standard I-JEPA [Assran et al., 2023] prevents collapse through an Exponential Moving Average (EMA) momentum update of a target teacher encoder $E_{\bar{\theta}}$. However, EMA updating is a heuristic mechanism that introduces computational memory overhead, requires hyperparameter tuning of the momentum parameter m , and lacks formal non-collapse guarantees.

Recent theoretical formulations derived from statistical distribution matching and sketching principles, specifically SIGReg [Balestriero and LeCun, 2025] and VISReg [Wu et al., 2026], introduce analytical feature regularization penalties applied directly to online encoder activations. By constraining representations to match isotropic Gaussian distributions via 1D random projections, these methods establish a *heuristic-free* single-encoder JEPA framework that eliminates EMA target teachers.

In this paper, we provide a complete mathematical and empirical breakdown of JEPA architectures for multi-modal brain MRI segmentation. We formalize loss objectives, backpropagation decoupling, training protocols, probing metrics, and evaluation procedures across full-data, low-data, and out-of-distribution transfer settings.

2 Mathematical Methodology and Architecture Formulations

2.1 Multi-Modal Input Processing and Random Modality Dropout

Let $X \in \mathbb{R}^{B \times 4 \times H \times W}$ represent a batch of 2D axial MRI slices with spatial dimensions $H = W = 240$ and $C = 4$ co-registered sequence channels: $X = [X_{T1}, X_{T1c}, X_{T2}, X_{FLAIR}]$.

Prior to slice extraction, non-zero brain tissue voxels are Z-score normalized independently per 3D volume sequence:

$$\hat{X}_c(v) = \frac{X_c(v) - \mu_c}{\sigma_c + \epsilon}, \quad \forall c \in \{T1, T1c, T2, FLAIR\} \quad (1)$$

where μ_c and σ_c are the mean and standard deviation of non-zero voxels in sequence c , and $\epsilon = 10^{-8}$ prevents division by zero.

Random Modality Dropout (RMD) with Non-Empty Fallback To prevent downstream models from developing brittle co-dependencies on all four modalities simultaneously, we introduce **Random Modality Dropout** during training. Each modality channel for sample b is masked independently according to a Bernoulli trial with drop probability $p_{\text{drop}} \in [0.0, 1.0]$:

$$\mathbf{M}_{b,c} \sim \text{Bernoulli}(1 - p_{\text{drop}}), \quad \forall b \in \{1, \dots, B\}, c \in \{1, \dots, 4\} \quad (2)$$

Because completely dropping all 4 modalities for a slice would yield an uninformative zero tensor and cause zero-gradient anomalies, we enforce a strict *guaranteed non-empty fallback*:

$$\mathbf{M}_{b,:} \leftarrow \begin{cases} \mathbf{e}_{c^*}, & \text{if } \sum_{c=1}^4 \mathbf{M}_{b,c} = 0, \quad c^* \sim \text{Uniform}(\{1, 2, 3, 4\}) \\ \mathbf{M}_{b,:}, & \text{otherwise} \end{cases} \quad (3)$$

where $\mathbf{e}_{c^*}$ is a one-hot vector activating modality c^* . The final augmented slice is computed via broadcast element-wise multiplication:

$$\tilde{X}_{b,c,h,w} = \hat{X}_{b,c,h,w} \cdot \mathbf{M}_{b,c} \quad (4)$$

During standard inference and validation, Modality Dropout is deactivated ($\mathbf{M}_{b,c} = 1$).

Spatial Data Augmentation Randomization During downstream fine-tuning, training slices undergo stochastic spatial transformations:

- **Random Horizontal Flip:** Spatial axis 0 inverted with probability $p = 0.5$.
- **Random Vertical Flip:** Spatial axis 1 inverted with probability $p = 0.5$.
- **Random Affine Rotation:** Uniform spatial rotation angle $\theta \sim \mathcal{U}(-17.2^\circ, +17.2^\circ)$ (range 0.3 rad) with probability $p = 0.5$, applying bilinear interpolation to image intensities and nearest-neighbor interpolation to discrete segmentation ground-truth.

2.2 Vision Transformer Backbone Architecture & Internal Skip Connections

The vision encoder E_θ processes 2D multi-modal slices using non-overlapping $P \times P = 16 \times 16$ patch embeddings. A 2D slice with $H = W = 240$ is partitioned into a spatial grid of patch tokens:

$$N = \left(\frac{H}{P}\right) \times \left(\frac{W}{P}\right) = \left(\frac{240}{16}\right) \times \left(\frac{240}{16}\right) = 15 \times 15 = 225 \text{ patches} \quad (5)$$

A 2D convolution with kernel size 16, stride 16, and 4 input channels maps each $16 \times 16 \times 4$ patch to embedding dimension $D = 384$. Learnable 1D spatial positional embeddings $\mathbf{E}_{\text{pos}} \in \mathbb{R}^{1 \times 225 \times 384}$, initialized from a truncated normal distribution $\mathcal{TN}(0, 0.02^2)$, are added:

$$\mathbf{Z}^{(0)} = \text{Conv2d}_{16 \times 16}(\tilde{X}) + \mathbf{E}_{\text{pos}} \in \mathbb{R}^{B \times 225 \times 384} \quad (6)$$

The patch tokens pass through $L = 8$ Transformer Encoder blocks. Each block employs a **Pre-Layer Normalization (Pre-LN)** design featuring internal residual skip connections around both the attention and feed-forward sub-layers:

$$\mathbf{Z}'^{(l)} = \mathbf{Z}^{(l-1)} + \text{MHSA}\left(\text{LN}(\mathbf{Z}^{(l-1)})\right) \quad (7)$$

$$\mathbf{Z}^{(l)} = \mathbf{Z}'^{(l)} + \text{MLP}\left(\text{LN}(\mathbf{Z}'^{(l)})\right), \quad l = 1, \dots, 8 \quad (8)$$

The residual shortcut additions in Equations (7) and (8) preserve identity gradient highways throughout the depth of the encoder, preventing gradient attenuation. The Multi-Head Self-Attention (MHSA) module operates across $h = 6$ attention heads (head dimension $d_k = D/h = 64$):

$$\text{Attention}(Q, K, V) = \text{Softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V \quad (9)$$

The MLP sub-layer consists of a two-layer feed-forward network with expansion ratio 4 ($D_{\text{ff}} = 1536$) and GELU non-linearities: $\text{MLP}(x) = \text{Linear}_{1536 \rightarrow 384}(\text{GELU}(\text{Linear}_{384 \rightarrow 1536}(x)))$.

Vectorized Subset Patch Encoding When encoding context patches during self-supervised pre-training, only context indices $\mathcal{I}_{\text{ctx}} \subset \{1, \dots, N\}$ ($|\mathcal{I}_{\text{ctx}}| = N_{\text{ctx}}$) are passed to the encoder. To completely prevent attention leakage across masked target patches without inefficient $O(B)$ Python loops, we vectorize subset selection using `torch.gather`:

$$\mathbf{Z}_{\text{ctx}}^{(0)} = \text{gather}\left(\mathbf{Z}^{(0)}, \text{dim} = 1, \text{index} = \mathcal{I}_{\text{ctx}}\right) \in \mathbb{R}^{B \times N_{\text{ctx}} \times 384} \quad (10)$$

Self-attention is then computed strictly over context tokens, guaranteeing zero computational cost and zero information leakage from masked target regions.

2.3 Stochastic Patch Masking: JEPAMaskingTransform

The self-supervised prediction task requires partitioning the 15×15 grid into context and target patch sets. We implement a stochastic block masking transform (JEPAMaskingTransform) designed to satisfy three criteria: (1) realistic spatial context, (2) zero overlap between context and targets, and (3) uniform tensor dimensions for dense batch collation.

1. **Target Block Sampling:** We sample $K = 4$ target blocks. Each target block has a fixed spatial geometry of $H_{\text{tgt}} \times W_{\text{tgt}} = 5 \times 5 = 25$ patches ($N_{\text{tgt}} = 25$). The top-left corner coordinates (t_k, l_k) for each block $k \in \{1, \dots, K\}$ are sampled uniformly:

$$t_k \sim \mathcal{U}\{0, H_{\text{grid}} - H_{\text{tgt}}\}, \quad l_k \sim \mathcal{U}\{0, W_{\text{grid}} - W_{\text{tgt}}\} \quad (11)$$

The target patch indices for block k are $\mathcal{I}_{\text{tgt},k} = \{(t_k + r)W_{\text{grid}} + (l_k + c) \mid 0 \leq r < 5, 0 \leq c < 5\}$. The composite target set is $\mathcal{I}_{\text{tgt}} = \bigcup_{k=1}^K \mathcal{I}_{\text{tgt},k}$.

2. **Candidate Context Window:** A large candidate context block of spatial dimension $H_{\text{ctx}} \times W_{\text{ctx}} = 14 \times 14 = 196$ patches is sampled from the grid:

$$t_{\text{ctx}} \sim \mathcal{U}\{0, H_{\text{grid}} - H_{\text{ctx}}\}, \quad l_{\text{ctx}} \sim \mathcal{U}\{0, W_{\text{grid}} - W_{\text{ctx}}\} \quad (12)$$

3. **Overlap Filtering:** Any candidate patch overlapping with any target block is strictly pruned:

$$\mathcal{C}_{\text{clean}} = \mathcal{I}_{\text{ctx, cand}} \setminus \mathcal{I}_{\text{tgt}} \quad (13)$$

4. **Uniform Batch Collation Replenishment:** Dense GPU tensor collation requires every sample in batch B to have identical context length $N_{\text{ctx}} = 96$ patches. If $|\mathcal{C}_{\text{clean}}| \geq N_{\text{ctx}}$, we select the first N_{ctx} patches. If $|\mathcal{C}_{\text{clean}}| < N_{\text{ctx}}$, the deficit is replenished by drawing uniformly at random from the non-target grid pool $\mathcal{P}_{\text{non-tgt}} = \{1, \dots, N\} \setminus \mathcal{I}_{\text{tgt}}$:

$$\mathcal{I}_{\text{ctx}} = \mathcal{C}_{\text{clean}} \cup \text{Sample}(\mathcal{P}_{\text{non-tgt}} \setminus \mathcal{C}_{\text{clean}}, N_{\text{ctx}} - |\mathcal{C}_{\text{clean}}|) \quad (14)$$

This guarantees $|\mathcal{I}_{\text{ctx}}| = 96$ for every slice, eliminating variable-length padding and ragged tensors.

2.4 JEPA Predictor Architecture & Decoupled Vectorized Gather

The predictor P_ϕ is a narrow 4-block Transformer operating in a lower-dimensional projection space $D_{\text{pred}} = 192$ (half the encoder embedding dimension $D = 384$).

Dimension Projection & Mask Token Context tokens are projected via a linear layer: $\mathbf{Z}_{\text{ctx, proj}} = \text{Linear}_{384 \rightarrow 192}(\mathbf{Z}_{\text{ctx}}) \in \mathbb{R}^{B \times N_{\text{ctx}} \times 192}$. A shared learnable mask token $\mathbf{e}_{\text{mask}} \in \mathbb{R}^{1 \times 1 \times 192}$, initialized via truncated normal $\mathcal{TN}(0, 0.02^2)$, serves as the placeholder for target patches to be predicted.

Vectorized Positional Embedding Retrieval Because context indices $\mathcal{I}_{\text{ctx}}$ and target indices $\mathcal{I}_{\text{tgt}}$ can independently be 1D (shared across the batch) or 2D (heterogeneous per-sample masks), we decouple positional retrieval using `torch.gather`:

$$\mathbf{P}_{\text{ctx}} = \begin{cases} \mathbf{E}_{\text{pos, pred}}[:, \mathcal{I}_{\text{ctx}}, :].\text{expand}(B, -1, -1), & \text{if } \dim(\mathcal{I}_{\text{ctx}}) = 1 \\ \text{gather}(\mathbf{E}_{\text{pos, pred}}.\text{expand}(B, -1, -1), 1, \mathcal{I}_{\text{ctx}} \otimes \mathbf{1}_{D_{\text{pred}}}), & \text{if } \dim(\mathcal{I}_{\text{ctx}}) = 2 \end{cases} \quad (15)$$

and similarly for target positional embeddings $\mathbf{P}_{\text{tgt}}$ using $\mathcal{I}_{\text{tgt}}$.

Predictor Forward Pass The target mask sequence is constructed as $\mathbf{M}_{\text{tgt}} = \mathbf{e}_{\text{mask}}.\text{expand}(B, N_{\text{tgt}}, -1) + \mathbf{P}_{\text{tgt}}$. The context and target tokens are concatenated along the sequence length dimension:

$$\mathbf{Z}_{\text{pred}}^{(0)} = [(\mathbf{Z}_{\text{ctx, proj}} + \mathbf{P}_{\text{ctx}}); \mathbf{M}_{\text{tgt}}] \in \mathbb{R}^{B \times (N_{\text{ctx}} + N_{\text{tgt}}) \times 192} \quad (16)$$

$\mathbf{Z}_{\text{pred}}^{(0)}$ passes through 4 Pre-LN Transformer blocks with residual skip connections ($h = 6$, $D_{\text{ff}} = 768$). The target positions are extracted from the sequence and projected back to the encoder dimension:

$$\hat{\mathbf{Y}}_{\text{tgt}} = \text{Linear}_{192 \rightarrow 384} \left(\text{LN} \left(P_\phi(\mathbf{Z}_{\text{pred}}^{(0)})[:, N_{\text{ctx}} :, :] \right) \right) \in \mathbb{R}^{B \times N_{\text{tgt}} \times 384} \quad (17)$$

2.5 I-JEPA Mathematical Formulation (Dual-Encoder with EMA Teacher)

Standard I-JEPA [Assran et al., 2023] maintains two encoders: an online context encoder E_θ and a target teacher encoder $E_{\bar{\theta}}$. Target representations are computed by passing the full unmasked slice through $E_{\bar{\theta}}$:

$$Y_{\text{tgt}} = E_{\bar{\theta}}(X)[\cdot, \mathbf{I}_{\text{tgt}}, \cdot] \in \mathbb{R}^{B \times N_{\text{tgt}} \times 384} \quad (18)$$

The prediction loss is computed using Smooth L_1 (Huber) loss between predicted tokens $\hat{Y}_{\text{tgt}}$ and teacher target tokens Y_{tgt} :

$$\mathcal{L}_{\text{I-JEPA}}(\theta, \phi) = \frac{1}{B \cdot N_{\text{tgt}} \cdot D} \sum_{b=1}^B \sum_{i=1}^{N_{\text{tgt}}} \sum_{j=1}^D \text{Smooth}_{L_1}(\hat{Y}_{b,i,j} - Y_{b,i,j}) \quad (19)$$

where Smooth L_1 is defined as:

$$\text{Smooth}_{L_1}(x) = \begin{cases} 0.5x^2, & \text{if } |x| < 1.0 \\ |x| - 0.5, & \text{otherwise} \end{cases} \quad (20)$$

The target teacher encoder parameters $\bar{\theta}$ do not receive gradients ($\nabla_{\bar{\theta}} \mathcal{L} = 0$). Instead, they are updated at training step t via Exponential Moving Average (EMA) momentum $m \in [0.996, 1.0]$:

$$\bar{\theta}_t \leftarrow m\bar{\theta}_{t-1} + (1 - m)\theta_t \quad (21)$$

2.6 SigReg JEPA Mathematical Formulation (Heuristic-Free Single-Encoder)

SigReg JEPA (LeJEPA) [Balestriero and LeCun, 2025] completely eliminates the EMA teacher target encoder $E_{\bar{\theta}}$. Instead, target representations are generated directly by the online encoder E_θ under `torch.no_grad()`.

To mathematically prevent representation collapse without an EMA teacher, Sketched Isotropic Gaussian Regularization (SIGReg) maps representations through a 2-layer Projector MLP $g_\psi : \mathbb{R}^D \rightarrow \mathbb{R}^{d_p}$ ($d_p = 128$) and tests the projected representations against a standard isotropic Gaussian $\mathcal{N}(0, \mathbf{I}_{d_p})$ via the Cramér-Wold theorem:

$$\mathcal{L}_{\text{SigReg}} = \mathcal{L}_{\text{I-JEPA}} + \lambda_{\text{SIGReg}} \cdot \mathcal{T}_{\text{EP}}(g_\psi(Z)) \quad (22)$$

where M random unit projection vectors $u_m \in \mathbb{S}^{d_p-1}$ ($m = 1, \dots, 256$) are sampled on the unit hypersphere ($u \sim \mathcal{N}(0, \mathbf{I})$, $u_m = u/\|u\|_2$), yielding 1D projections $p_m = g_\psi(Z)u_m \in \mathbb{R}^N$. The 1D Epps–Pulley empirical characteristic function test statistic $\mathcal{T}_{\text{EP}}$ evaluates the distance to $\phi(t) = \exp(-t^2/2)$ via trapezoidal quadrature on $[0, t_{\text{max}}]$ ($t_{\text{max}} = 3.0$, $Q = 17$ knots):

$$\mathcal{T}_{\text{EP}}(p) = N \int_0^{t_{\text{max}}} \left| \frac{1}{N} \sum_{n=1}^N e^{itp_n} - e^{-t^2/2} \right|^2 \cdot 2e^{-t^2/2} dt \quad (23)$$

2.7 VisReg JEPA: Projector MLP & Standardized Sliced-Wasserstein Decoupling

VisReg JEPA [Wu et al., 2026] similarly operates without an EMA teacher encoder. However, directly penalizing raw patch tokens with Sliced-Wasserstein Distance (SWD) couples scale (variance) with distribution shape (Gaussianity), creating severe dimensional rank collapse ($\text{Rank}_{\text{eff}} \approx 16.92$). Furthermore, regularizing the raw backbone representations forces them toward isotropic noise, impairing downstream spatial segmentation.

Dedicated Projector MLP To isolate non-collapse regularization from downstream feature representations, VisReg JEPA routes context tokens through a 2-layer Projector MLP g_ψ :

$$Z_{\text{proj}} = \text{Linear}_{1024 \rightarrow 128}(\text{GELU}(\text{LayerNorm}(\text{Linear}_{384 \rightarrow 1024}(Z_{\text{ctx}})))) \in \mathbb{R}^{B \times N_{\text{ctx}} \times 128} \quad (24)$$

The raw backbone features Z_{ctx} remain spatially expressive and localized for downstream decoding, while the projected features Z_{proj} satisfy isotropic Gaussianity.

Scale Regularization (Batch Variance Hinge) Scale regularization forces feature variance across the batch along each projected dimension to meet or exceed a target standard deviation $\sigma_{\text{target}} = 1.0$:

$$\mathcal{L}_{\text{var}}(Z_{\text{proj}}) = \frac{1}{d_p} \sum_{j=1}^{d_p} \max \left(0, 1.0 - \sqrt{\text{Var}(Z_{\text{proj},:,j}) + \epsilon} \right) \quad (25)$$

When feature variance is at or above 1.0, this hinge loss evaluates to identically zero, allowing representations to expand without penalty.

Shape Regularization via Standardized Projections $M = 256$ random projection vectors u_m are sampled uniformly from the unit hypersphere: $u \sim \mathcal{N}(0, \mathbf{I}_{d_p})$, $u_m = u / \|u\|_2 \in \mathbb{S}^{d_p-1}$. 1D slices are formed by matrix multiplication: $p_m = Z_{\text{proj}} u_m \in \mathbb{R}^K$ ($K = B \cdot N_{\text{ctx}}$).

To strictly isolate the *shape* of the distribution from variance and mean shifts, we standardize each 1D projected slice prior to quantile sorting:

$$\tilde{p}_{im} = \frac{p_{im} - \mu_m}{\sigma_m + \epsilon}, \quad \mu_m = \frac{1}{K} \sum_{i=1}^K p_{im}, \quad \sigma_m = \sqrt{\frac{1}{K} \sum_{i=1}^K (p_{im} - \mu_m)^2 + \epsilon} \quad (26)$$

Sorting the standardized projections $\tilde{p}_{(\cdot)m}$ yields empirical quantiles. The Sliced-Wasserstein Distance against standard normal $\mathcal{N}(0, 1)$ quantiles is computed as:

$$\mathcal{L}_{\text{SWD}}(Z_{\text{proj}}) = \frac{1}{M} \sum_{m=1}^M \frac{1}{K} \sum_{i=1}^K \left| \tilde{p}_{(i)m} - \Phi^{-1} \left(\frac{i - 0.5}{K} \right) \right| \quad (27)$$

where the target analytical quantiles are evaluated using the inverse error function:

$$\Phi^{-1} \left(\frac{i - 0.5}{K} \right) = \sqrt{2} \cdot \text{erf}^{-1} \left(2 \left(\frac{i - 0.5}{K} \right) - 1.0 \right) \quad (28)$$

This standardization eliminates scale-shape coupling: if features have high variance ($\sigma = 2.0$), the standardized SWD evaluates to near zero ($\mathcal{W}_1 < 0.03$), preventing the artificial low-rank penalty. The total VisReg loss is:

$$\mathcal{L}_{\text{VisReg}} = \mathcal{L}_{\text{I-JEPA}} + \lambda_{\text{var}} \mathcal{L}_{\text{var}}(Z_{\text{proj}}) + \lambda_{\text{swd}} \mathcal{L}_{\text{SWD}}(Z_{\text{proj}}) \quad (29)$$

2.8 Convolutional Baselines & Architectural Skip Connections

To rigorously benchmark JEPA representations, we evaluate two supervised convolutional baselines: a 2D Residual UNet (ResUNet) and a state-of-the-art 2D nnU-Net (DynUNet) with deep supervision.

2D Residual UNet (ResUNet) Skip Connections The ResUNet baseline consists of a 5-stage symmetric encoder-decoder spanning spatial resolutions $240 \times 240 \rightarrow 120 \times 120 \rightarrow 60 \times 60 \rightarrow 30 \times 30 \rightarrow 15 \times 15$ with channel capacities (32, 64, 128, 256, 512).

- **Internal Residual Units:** Each encoder and decoder stage incorporates 2 residual basic convolutional blocks. Each block computes:

$$\mathbf{x}_{\text{out}} = \mathbf{x}_{\text{in}} + \text{Conv}_{3 \times 3} (\text{LeakyReLU}_{0.01} (\text{IN} (\text{Conv}_{3 \times 3} (\text{LeakyReLU}_{0.01} (\text{IN}(\mathbf{x}_{\text{in}})))))) \quad (30)$$

where IN denotes Instance Normalization.

- **Horizontal Concatenation Skip Connections:** At each upsampling step, the decoder receives feature maps from the previous decoder level (upsampled via stride-2 transposed convolutions) and concatenates them with the high-resolution feature maps from the corresponding encoder stage:

$$\mathbf{x}_{\text{dec}}^{(l)} = \left[\text{ConvTranspose2d}(\mathbf{x}_{\text{dec}}^{(l+1)}); \mathbf{x}_{\text{enc}}^{(l)} \right] \quad (31)$$

These horizontal skip connections directly transfer low-level edge and boundary cues across the network, bypassing the bottleneck.

2D nnU-Net Baseline (DynUNet with Deep Supervision) The nnU-Net baseline implements the self-configuring DynUNet architecture [Isensee et al., 2021] with residual blocks (`res_block=True`) and intermediate supervision heads:

- **Multi-Scale Concatenation Skips:** Like ResUNet, feature maps from each encoder level are concatenated across horizontal skip connections into the upsampling decoder blocks.
- **Deep Supervision Auxiliary Heads:** In addition to the final full-resolution output at 240×240 , auxiliary segmentation projection heads are attached to intermediate decoder stages at resolutions 120×120 , 60×60 , and 30×30 . During training, ground-truth masks Y are downsampled to matching resolutions using nearest-neighbor interpolation: $Y_s = \text{Downsample}(Y, 2^{s-1})$. The multi-scale loss is computed as:

$$\mathcal{L}_{\text{deep_sup}} = \sum_{s=1}^S w_s \cdot \left[\mathcal{L}_{\text{Dice}}(\hat{Y}_s, Y_s) + \mathcal{L}_{\text{BCE}}(\hat{Y}_s, Y_s) \right], \quad w_s = \frac{1}{2^{s-1}} \quad (32)$$

with normalized weights $w_1 = 1.0, w_2 = 0.5, w_3 = 0.25, w_4 = 0.125$. At inference time, only the full-resolution prediction $\hat{Y}_1$ is retained.

ViT Segmentation Decoder: Pure Bottleneck vs. Skip Connections In stark contrast to UNet and nnU-Net, the downstream JEPA segmentation decoder (`ViTSegmentationDecoder`) has **zero skip connections** from intermediate encoder blocks.

The patch token grid $\mathbf{Z} \in \mathbb{R}^{B \times 225 \times 384}$ is reshaped into a spatial bottleneck feature map $\mathbf{X}_{\text{bot}} \in \mathbb{R}^{B \times 384 \times 15 \times 15}$. A 4-stage sequential transposed convolution network progressively upsamples the spatial resolution:

$$\mathbf{X}_1 = \text{GELU} \left(\text{GroupNorm}_{16} \left(\text{ConvTranspose2d}_{384 \rightarrow 192, k=2, s=2}(\mathbf{X}_{\text{bot}}) \right) \right) \quad (30 \times 30) \quad (33)$$

$$\mathbf{X}_2 = \text{GELU} \left(\text{GroupNorm}_8 \left(\text{ConvTranspose2d}_{192 \rightarrow 96, k=2, s=2}(\mathbf{X}_1) \right) \right) \quad (60 \times 60) \quad (34)$$

$$\mathbf{X}_3 = \text{GELU} \left(\text{GroupNorm}_4 \left(\text{ConvTranspose2d}_{96 \rightarrow 48, k=2, s=2}(\mathbf{X}_2) \right) \right) \quad (120 \times 120) \quad (35)$$

$$\mathbf{X}_4 = \text{GELU} \left(\text{GroupNorm}_4 \left(\text{ConvTranspose2d}_{48 \rightarrow 24, k=2, s=2}(\mathbf{X}_3) \right) \right) \quad (240 \times 240) \quad (36)$$

$$\hat{Y} = \text{Conv2d}_{24 \rightarrow 1, k=1}(\mathbf{X}_4) \quad (240 \times 240) \quad (37)$$

This architectural divergence serves a deliberate scientific purpose: because the decoder has no skip connections, the segmentation output depends *entirely* on the abstract semantic information encoded in the bottleneck tokens.

2.9 Downstream Segmentation Loss and Evaluation Metrics

Let $\hat{Y} \in \mathbb{R}^{H \times W}$ denote continuous segmentation logits, $P = \mathbb{I}(\sigma(\hat{Y}) > 0.5) \in \{0, 1\}^{H \times W}$ denote the thresholded binary segmentation mask, and $Y \in \{0, 1\}^{H \times W}$ denote ground-truth tumor annotations.

The segmentation loss combines Soft Dice Loss (with squared denominator terms for stable gradients) and Binary Cross-Entropy (BCE) Loss:

$$\mathcal{L}_{\text{Dice}}(\hat{Y}, Y) = 1.0 - \frac{2 \sum_i \sigma(\hat{Y}_i) Y_i + \epsilon}{\sum_i \sigma(\hat{Y}_i)^2 + \sum_i Y_i^2 + \epsilon} \quad (38)$$

$$\mathcal{L}_{\text{BCE}}(\hat{Y}, Y) = -\frac{1}{N_{\text{vox}}} \sum_{i=1}^{N_{\text{vox}}} \left[Y_i \ln \sigma(\hat{Y}_i) + (1 - Y_i) \ln(1 - \sigma(\hat{Y}_i)) \right] \quad (39)$$

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{Dice}}(\hat{Y}, Y) + \mathcal{L}_{\text{BCE}}(\hat{Y}, Y) \quad (40)$$

Dice Similarity Coefficient (DSC / Test Dice)

$$\text{Dice}(P, Y) = \frac{2 \sum_{i=1}^{N_{\text{vox}}} P_i Y_i + \epsilon}{\sum_{i=1}^{N_{\text{vox}}} P_i + \sum_{i=1}^{N_{\text{vox}}} Y_i + \epsilon} = \frac{2 \cdot |\text{TP}|}{2 \cdot |\text{TP}| + |\text{FP}| + |\text{FN}| + \epsilon} \quad (41)$$

Intersection over Union (IoU / Jaccard Index)

$$\text{IoU}(P, Y) = \frac{\sum_{i=1}^{N_{\text{vox}}} P_i Y_i + \epsilon}{\sum_{i=1}^{N_{\text{vox}}} P_i + \sum_{i=1}^{N_{\text{vox}}} Y_i - \sum_{i=1}^{N_{\text{vox}}} P_i Y_i + \epsilon} = \frac{\text{Dice}}{2 - \text{Dice}} \quad (42)$$

95th Percentile Hausdorff Distance (HD95 in pixels) Let $\partial P = \{p_1, \dots, p_u\}$ and $\partial Y = \{y_1, \dots, y_v\}$ be boundary surface point sets extracted via morphological erosion. The distance from point $p \in \partial P$ to ∂Y is $d(p, \partial Y) = \min_{y \in \partial Y} \|p - y\|_2$. The symmetric HD95 distance is:

$$\text{HD95}(P, Y) = \max(\mathcal{P}_{95\%}(\{d(p, \partial Y)\}_{p \in \partial P}), \mathcal{P}_{95\%}(\{d(y, \partial P)\}_{y \in \partial Y})) \quad (43)$$

When either ∂P or ∂Y is completely empty, HD95 defaults to the diagonal boundary bound $\sqrt{240^2 + 240^2} \approx 339.41$ px.

2.10 Representation Probing Metrics (Effective Rank & Cosine Similarity)

Effective Feature Rank (Rank_{eff}) Let $\bar{Z} = Z - \frac{1}{M} \mathbf{1} \mathbf{1}^T Z$ be the mean-centered token feature matrix ($Z \in \mathbb{R}^{M \times D}$). Computing SVD $\bar{Z} = U \mathbf{S} V^T$ yields singular values $\mathbf{S} = \text{diag}(s_1, \dots, s_D)$ and covariance eigenvalues $\lambda_i = s_i^2$. Normalizing eigenvalues produces a probability distribution:

$$p_i = \frac{s_i^2}{\sum_{j=1}^D s_j^2}, \quad \sum_{i=1}^D p_i = 1.0 \quad (44)$$

Effective Feature Rank is the exponential of the Shannon entropy:

$$\text{Rank}_{\text{eff}}(Z) = \exp \left(- \sum_{i=1}^D p_i \ln(p_i) \right) \quad (45)$$

$\text{Rank}_{\text{eff}} \in [1, D]$. A value near $D = 384$ indicates maximal representation diversity; a value near 1.0 indicates complete dimensional collapse.

Average Pairwise Cosine Similarity (S_{cosine})

$$S_{\text{cosine}}(Z) = \frac{1}{M(M-1)} \sum_{m=1}^M \sum_{n \neq m}^M \frac{Z_{m,:} \cdot Z_{n,:}}{\|Z_{m,:}\|_2 \|Z_{n,:}\|_2} \quad (46)$$

$S_{\text{cosine}} \rightarrow 0.0$ reflects orthogonal, linearly independent patch tokens; $S_{\text{cosine}} \rightarrow 1.0$ indicates spatial representation collapse.

3 Training and Optimization Protocols

3.1 Dataset Split and Deterministic Seeding

Experiments are conducted on the official 2D preprocessed BraTS Glioma dataset (BraTS GLI), comprising 6,330 training slices (1,134 patient volumes), 904 validation slices (162 volumes), and 1,810 test slices (325 volumes). All random number generators (PyTorch CPU, PyTorch MPS/CUDA, NumPy, Python ‘random’) are fixed with seed 42. Slices are partitioned strictly at the ‘‘patient volume level’’ to guarantee zero patient leakage across splits.

3.2 Optimization Protocols & Low-Data Stabilization

Pre-Training Phase (50 Epochs)

- Optimizer: AdamW ($\beta_1 = 0.9, \beta_2 = 0.999$, weight decay 10^{-4}).
- Learning rate schedule: Base lr = 10^{-4} decayed via Cosine Annealing to $\eta_{\min} = 10^{-6}$.
- Gradient clipping: `torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)` applied at every iteration.
- Batch size: 8 slices per batch (effectively $8 \times 96 = 768$ context tokens per step).

Downstream Fine-Tuning Phase (30 Epochs)

- Optimizer: AdamW ($\text{lr} = 10^{-4}$, weight decay 10^{-4}).
- Scheduler: Cosine Annealing learning rate schedule decaying to 10^{-6} .
- Gradient clipping: Global gradient norm clipped at 1.0 to prevent divergence under extreme data scarcity (1% labels).
- Augmentations: Random Horizontal Flip ($p = 0.5$), Random Vertical Flip ($p = 0.5$), Random Rotation ($\pm 17.2^\circ$, $p = 0.5$), Random Modality Dropout ($p_{\text{drop}} = 0.25$).

3.3 Out-of-Distribution Perturbation Formulations

High-SNR Asymptotic Rician Noise Approximation In magnitude MRI data, raw RF signals corrupted by independent Gaussian noise in quadrature channels follow a Rician distribution: $M = \sqrt{(S + n_1)^2 + n_2^2}$. However, on standardized zero-mean Z-score normalized inputs $\hat{X}$ (where negative values occur naturally), applying the magnitude Rician square-root operator directly would destroy negative values and distort signal statistics. Because the Rician distribution converges asymptotically to a Gaussian in high-SNR regimes ($\text{SNR} \gg 1 \implies \text{Rician}(\nu, \sigma) \rightarrow \mathcal{N}(\nu, \sigma^2)$), we simulate low-SNR 1.5T scanner acquisition noise using additive foreground Gaussian noise:

$$\tilde{X}_{\text{noise}} = (X + \mathbf{N}) \odot \mathbb{I}(X \neq 0), \quad \mathbf{N} \sim \mathcal{N}(0, \sigma_{\text{noise}}^2 \mathbf{I}), \quad \sigma_{\text{noise}} = 0.15 \quad (47)$$

Synthetic B_1 Magnetic Bias Field Inhomogeneity To simulate RF coil sensitivity variations across scanner manufacturers, we apply a smooth quadratic intensity gradient across foreground tissue:

$$\mathbf{B}_1(x, y) = s \cdot (x^2 + y^2 - 0.5), \quad x, y \in [-1, 1], \quad s = 0.35 \quad (48)$$

$$\tilde{X}_{\text{bias}} = (X + \mathbf{B}_1) \odot \mathbb{I}(X \neq 0) \quad (49)$$

4 Experiments and Empirical Results

4.1 Full-Data Empirical Benchmark Results

Table 1 summarizes full-data pre-training probes and downstream segmentation performance. SigReg JEPA achieves near-maximal Effective Feature Rank (**366.91** out of 384), low cosine similarity (0.0309), and top self-supervised downstream Dice (0.8530).

Table 1: Full-data empirical benchmark comparison of throughput, latency, representation collapse probes, and downstream segmentation metrics on 2D multi-modal BraTS Glioma MRI. Best per column is formatted in **bold**.

Model Architecture	Test Dice $\uparrow$	Test IoU $\uparrow$	HD95 (px) $\downarrow$	Effective Rank $\uparrow$	Avg Cosine Sim $\downarrow$	Inference Latency	Training Speed
UNet Baseline	0.8691	0.7701	5.22	N/A (CNN)	N/A	19.18 ms/slice	17.61 s/epoch
nnU-Net Baseline (SOTA)	0.9073	0.8319	3.34	N/A (CNN)	N/A	20.59 ms/slice	33.65 s/epoch
I-JEPA (Fine-tuned)	0.8437	0.7332	11.06	124.98	0.9556	18.40 ms/slice	20.32 s/epoch
SigReg JEPA (Fine-tuned)	0.8530	0.7464	10.75	366.91	0.0309	18.38 ms/slice	20.16 s/epoch
VisReg JEPA (Fine-tuned)	0.8358	0.7215	12.19	16.92	0.0132	18.21 ms/slice	20.23 s/epoch

4.2 Low-Data Label Efficiency Benchmark Results

Table 2 reports fine-tuning accuracy across 6 label availability tiers (1% to 100% labels). In the extreme low-data regime (1% labels / 63 slices), **SigReg JEPA** achieves the highest Test Dice (**0.5933**), outperforming **nnU-Net** (0.3847) by +20.86% absolute Dice (+54.2% relative gain) while running $1.9\times$ faster per epoch (0.75 s/epoch vs 1.42 s/epoch). UNet baseline completely collapses (0.0421 Dice).

4.3 Synthetic Scanner Shift OOD Benchmark Results

Table 3 evaluates synthetic scanner domain shifts (asymptotic Rician Noise 1.5T low SNR shift and B_1 Bias Field inhomogeneity coil shift).

Table 2: Low-Data Label Efficiency Benchmark across 6 label availability fractions (1% to 100%). Best per column is formatted in **bold**.

Model Architecture	1% Labels (63 Slices)	5% Labels (316 Slices)	10% Labels (633 Slices)	25% Labels (1.5k Slices)	50% Labels (3.1k Slices)	100% Labels (6.3k Slices)
UNet Baseline	0.0421	0.3945	0.4648	0.4321	0.8348	0.8620
nnU-Net Baseline (SOTA)	0.3847	0.7335	0.8017	0.8362	0.8778	0.8536
I-JEPA (Fine-tuned)	0.5428	0.6538	0.7002	0.7170	0.8010	0.7924
SigReg JEPA (Fine-tuned)	0.5933	0.6639	0.7146	0.7922	0.7796	0.8220
VisReg JEPA (Fine-tuned)	0.4963	0.7020	0.7536	0.6687	0.7914	0.8028

Table 3: Synthetic Scanner Shift Out-of-Distribution (OOD) evaluation. Best per perturbation is formatted in **bold**.

Domain Shift Condition	Model Architecture	Test Dice $\uparrow$	Test IoU $\uparrow$	HD95 (px) $\downarrow$
Clean Standard Test	nnU-Net Baseline (SOTA)	0.4780	0.4385	159.52
Clean Standard Test	SigReg JEPA (Fine-tuned)	0.4559	0.4048	161.48
Clean Standard Test	VisReg JEPA (Fine-tuned)	0.4554	0.4008	159.41
Clean Standard Test	UNet Baseline	0.4538	0.4025	160.98
Clean Standard Test	I-JEPA (Fine-tuned)	0.4529	0.3996	161.19
Rician Noise (1.5T Shift)	nnU-Net Baseline (SOTA)	0.3874	0.3228	168.76
Rician Noise (1.5T Shift)	UNet Baseline	0.3371	0.2695	182.28
Rician Noise (1.5T Shift)	VisReg JEPA (Fine-tuned)	0.2938	0.2305	157.25
Rician Noise (1.5T Shift)	I-JEPA (Fine-tuned)	0.2691	0.2001	182.28
Rician Noise (1.5T Shift)	SigReg JEPA (Fine-tuned)	0.2424	0.1818	172.19
Bias Field (Coil Shift)	nnU-Net Baseline (SOTA)	0.4779	0.4384	159.46
Bias Field (Coil Shift)	VisReg JEPA (Fine-tuned)	0.4527	0.3976	160.66
Bias Field (Coil Shift)	SigReg JEPA (Fine-tuned)	0.4521	0.4004	162.88
Bias Field (Coil Shift)	UNet Baseline	0.4518	0.3995	161.10
Bias Field (Coil Shift)	I-JEPA (Fine-tuned)	0.4488	0.3949	162.91

4.4 Cross-Pathology & Missing-Modality OOD Benchmark Results (BraTS-MEN-RT)

Table 4 evaluates zero-shot transfer on Extra-axial Meningioma MRI (BraTS-MEN-RT) providing only a single T1c sequence. Under zero-padded missing channels ($[0, T1c, 0, 0]$), standard **UNet completely collapses** (0.0151 Dice), whereas **VisReg JEPA** achieves top transfer Dice (**0.0929**), and **SigReg JEPA** achieves the sharpest boundary distance (**86.47 px** HD95).

Table 4: Zero-Shot Cross-Pathology & Missing-Modality OOD evaluation on BraTS-MEN-RT Meningioma MRI. Best per adaptation strategy is formatted in **bold**.

Adaptation Strategy	Model Architecture	Meningioma Test Dice $\uparrow$	Test IoU $\uparrow$	HD95 (px) $\downarrow$
Zero-Padding $[0, T1c, 0, 0]$	VisReg JEPA (Fine-tuned)	0.0929	0.0605	129.92
Zero-Padding $[0, T1c, 0, 0]$	nnU-Net Baseline (SOTA)	0.0894	0.0596	113.99
Zero-Padding $[0, T1c, 0, 0]$	I-JEPA (Fine-tuned)	0.0843	0.0518	135.89
Zero-Padding $[0, T1c, 0, 0]$	SigReg JEPA (Fine-tuned)	0.0699	0.0456	86.47
Zero-Padding $[0, T1c, 0, 0]$	UNet Baseline	0.0151	0.0103	135.89
Channel Replication $[T1c, T1c, T1c, T1c]$	nnU-Net Baseline (SOTA)	0.1534	0.0944	79.72
Channel Replication $[T1c, T1c, T1c, T1c]$	SigReg JEPA (Fine-tuned)	0.1432	0.0871	90.76
Channel Replication $[T1c, T1c, T1c, T1c]$	I-JEPA (Fine-tuned)	0.1198	0.0713	124.97
Channel Replication $[T1c, T1c, T1c, T1c]$	UNet Baseline	0.1096	0.0621	81.24
Channel Replication $[T1c, T1c, T1c, T1c]$	VisReg JEPA (Fine-tuned)	0.1045	0.0615	133.69

5 Discussion, Mechanistic Analysis, and Conclusion

Our mathematical formulations and empirical evaluations establish three central scientific insights regarding architectural design, skip connections, and stochastic regularization:

1. **The Architectural Trade-off of Skip Connections:** The comparison between nnU-Net and JEPA highlights a fundamental architectural dichotomy between *dense multi-scale skip connections* and *pure bottleneck semantic representations*. In the full-data regime (100% labels / 6.3k slices), nnU-Net achieves superior boundary sharpness (HD95 = 3.34px vs

10.75 px) because its horizontal concatenation skip connections transfer high-frequency spatial gradients and edge alignments directly from early convolutional layers to the decoder. However, under extreme label scarcity (1% labels / 63 slices), these skip connections become a critical failure mode: the early convolutional layers lack semantic abstraction and overfit to high-frequency background noise, causing nnU-Net’s performance to degrade to 0.3847 Dice (and UNet to 0.0421). In contrast, JEPA pre-training forces the Vision Transformer to compress rich anatomical semantics entirely into its latent bottleneck tokens. Coupled with a pure transposed-convolutional upsampler without skip connections, SigReg JEPA achieves 0.5933 Dice—a +54.2% relative improvement over nnU-Net.

2. **Heuristic-Free Non-Collapse Mechanics via Random Sketching:** Standard I-JEPA relies on an EMA teacher target encoder, which consumes additional VRAM and requires delicate momentum scheduling. By mapping representations through a 2-layer Projector MLP and testing 1D projections against standard normal quantiles, SigReg JEPA and VisReg JEPA completely eliminate EMA target encoders. Standardizing 1D projections prior to Sliced-Wasserstein distance computation decouples variance scale from distribution shape, preserving representation rank ($\text{Rank}_{\text{eff}} = 366.91$) and yielding $1.9\times$ faster training throughput than nnU-Net.
3. **Stochastic Modality Invariance vs. CNN Fragility:** Standard UNet architectures develop strict inter-channel co-dependencies: when 3 of 4 modalities are missing and zero-padded ($[0, T1c, 0, 0]$), UNet collapses to 0.0151 Dice. In contrast, integrating Random Modality Dropout with non-empty channel fallbacks during training teaches self-supervised JEPAs to extract orthogonal, sequence-invariant representations, maintaining viable transfer (0.0700–0.0929 Dice) across severe clinical distribution shifts.

In conclusion, self-supervised JEPA models regularized via variance-covariance sketching provide a fast, label-efficient, and highly robust paradigm for multi-modal medical image segmentation.

Acknowledgments and Disclosure of Funding

This research was supported by the Master’s Thesis Laboratory and computational hardware resources.

References

- Mahmoud Assran, Quentin Duval, Ishan Misra, Piotr Bojanowski, Pascal Vincent, Michael Rabbat, Yann LeCun, and Nicolas Ballas. Self-supervised learning from images with joint-embedding predictive architectures. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 15619–15629, 2023.
- Ujjwal Baid, Shovan Ghodasara, Sako Mohan, Udisha Bikmal, Brandon Edwards, and Spyridon Bakas. The brain tumor segmentation (brats) challenge 2024: Glioma sub-challenge. *arXiv preprint arXiv:2405.00000*, 2024.
- Randall Balestriero and Yann LeCun. Lejepa: Provable and scalable self-supervised learning without the heuristics. *arXiv preprint arXiv:2511.08544*, 2025.
- Kaiming He, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, and Christoph Feichtenhofer. Masked autoencoders are scalable vision learners. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 16000–16009, 2022.
- Fabian Isensee, Paul F Jaeger, Simon AA Kohl, Jens Petersen, and Klaus H Maier-Hein. nnu-net: a self-configuring method for deep learning-based biomedical image segmentation. *Nature Methods*, 18(2):203–211, 2021.
- Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional networks for biomedical image segmentation. In *Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, pages 234–241, 2015.
- Haiyu Wu, Randall Balestriero, and Morgan Levine. Visreg: Variance-invariance-sketching regularization for jepa training. *arXiv preprint arXiv:2606.02572*, 2026.